

Object Detection

Halil Eralp Koçuş

January 24, 2020

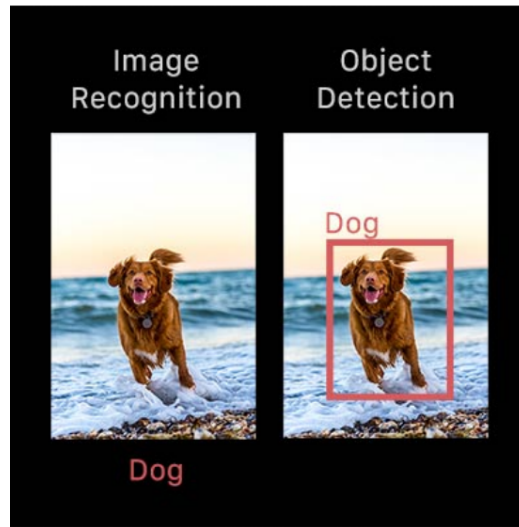


Figure 1: Difference between recognition and detection

1 Object Detection

Object detection is a concept of computer vision in which trained models detect objects with their category and location in given images or videos. As you can see in the figure 1, object detection is both classifying and localizing the object instances in given images.

1.1 Motivation

Various models are developed through years in the field of object detection. All these models are developed using different methods. Although all of these models are based on convolutional neural networks, their way of identifying objects in given images or frames are varies. Also, their backbone networks, scanning methods, multi-scale object detection, error functions are varies. These varieties affect the performance of models in different aspects such as small and large object detections, speed of the model, etc.

The aim in this article is to investigate, analyze, and compare the performances of various state-of-the-art object detectors through the time trained on video data. Then, promising detectors will be selected and these selected detectors will be analyzed on Video Object Detection dataset on both mean average precision and frame rate per second basis.

1.2 Literature in Static Object Detection

Static object detection refers to object detection in single image. In this context, object detectors do not use temporal information about objects in given image or frame sequences. Each given images are considered separately. The aim of investigating static object detection is to make a generalization from single frame to multiple frames. Understanding the pros and cons of each object detectors in various static object detection cases are important to make a generalization and creating a baseline for video object detection which requires more features to be considered for a well- performing detector.

TBC!!

1.3 Literature in Video Object Detection

Literature in Video Object Detection

2 Detector Features

As mentioned in Motivation section, in this article, the aim is to analyze object detectors based on various detector features. In the following five subsections, the following features will be examined:

- Backbone Networks
- Scanning Methods
- Multi-scale Handling
- Loss Functions
- Bells-and-whistles

These listed features are the essential features for detectors. When one studies the object detectors, one can see that major increases on the performances of object detectors are caused by changes on these features. These features will be analyzed and shown in the following subsections and section.

2.1 Backbone Networks

Backbone networks are the initial part of a detector's architecture. The given frame is processed for the first time in backbone networks. These backbone networks are implemented as Fully Convolutional Neural Networks. The used networks are the ones that are already proven to be perform well in images. The most frequent used ones are as following:

- VGG
- ResNet
- DarkNet
- Hourglass

The main function of these networks is to compute convolutional feature map over given frame, so that, these extracted feature maps are used to detect and localize objects in given frames.

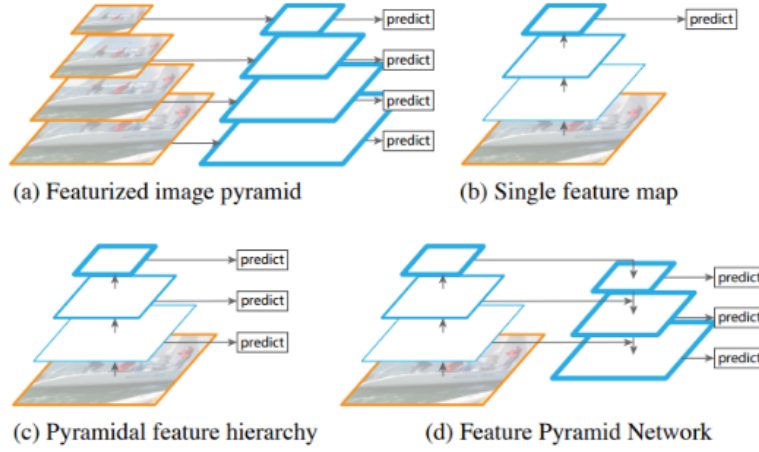


Figure 2: Different methods to use feature maps

2.2 Scanning Methods

There are many methods developed to scan given images through years. The following is a list for most frequent used scanning methods in object detectors:

- Sliding windows
- Region Proposals
- Grid Cells
- Anchor Boxes

All these methods have different advantages against each other. Also, Anchor Boxes are used together with other methods.

2.3 Multi-scale Handling

Multi-scale handling refers to detect objects that have different sizes in given frame. Multi-scale handling is a crucial feature for having a well performing detector since most of the objects in images have different sizes. There are various ways to handle this problem:

One of the solution (Fig. 2(a)) is to run detection multiple times that in each run, the resolution of input frame has to be changed. Then, all the detected objects has to be combined after all iterations are completed. In case of multiple detection for same object, a suppression has to be performed to reduce single detection. Although this method works well, its runtime is slow.

Another solution (Fig. 2(b)) is to use single feature map. In this method, single feature map is extracted from given frame and this feature map is passed through multiple convolutional layers to obtain a final feature map with more fine-grained features. Then, this feature map is used to predict the objects in given frame. This method is used to obtain fast detection but its performance is relatively worse than other methods.

Another solution (Fig. 2(c)) is to use pyramidal feature hierarchy in which multiple feature maps are used to make prediction. In this method, a feature map is extracted from given frame and as it is in second solution, this feature map is passed through multiple convolutional layers. However, in each layer, the extracted feature map is used to make prediction. This method is relatively works slower than second solution but it performs better.

Another solution (Fig. 2(d)) is to use feature pyramid network but this method will be examined detaily in section 3.2.7.

2.4 Loss Functions

Loss functions are used to measure the difference between ground-truth values and the predicted values in machine learning problems. Loss functions can also be called error functions. Thus, the aim is to minimize the value of loss function during training to have a better detector. In object detection problems, loss functions are composed of two parts: regression difference and classification difference. Loss functions are the measure of difference between predicted regression box and ground-truth regression box and also, the difference between the correct class for object and the predicted class for object. These functions can vary based on design choices and different loss functions have different advantages and disadvantages. The most frequent used loss functions as following:

- L1 Loss Function

L1 loss function minimizes the absolute differences between the predicted value and ground-truth value.

$$\sum_{i=1}^n (|y_t - y_p|)$$

- L2 Loss Function

L2 loss function minimizes the squared differences between the predicted value and ground-truth value.

$$\sum_{i=1}^n (y_t - y_i)^2$$

- Smooth L1 Loss Function

- Cross-Entropy Loss Function

Since L1 loss function operates on the absolute values, it does not affected by outliers. Therefore, L1 is more robust to outliers. On the other hand, L2 loss function operates on the squared values, therefore, L2 is not robust to outliers since outliers will cause a huge error value. However, if one analyzes and see there is not many outliers in data, then, using L2 loss function leads to a better training.




Figure 3: Network of Faster R-CNN

2.5 Bells-And-Whistles

3 Static Detector Types

In this section, state-of-the-art object detectors will be examined. The following detectors are the best-performing detectors in their publication time. These detectors are divided into two subsection that are two-stage detectors and one-stage detectors. The difference between these one-stage and two-stage detectors is their way to process given frame to detect objects.

In two-stage detectors, detections process takes two stages that are region proposal and object detection. In first phase, model generates region of interest that are the regions in which objects can be found most likely. Then, in second phase, a classifier processes these regions to detect objects.

In one-stage detectors, detection process takes only one stage. Detector runs over a dense sampling of possible locations in given frame. This approach converges faster but its performance might be worse than two-stage detectors.

3.1 Two-Stage Detectors

3.1.1 Faster R-CNN

Faster R-CNN consists of a region proposal network and a detection network which is improved version of Fast R-CNN as you can see in figure 3. The main improvement in Faster R-CNN is the shared convolutional layers between the RPN and the detection network, so that, the cost of detection is drastically reduced.

The RPN provides a simultaneous prediction of object bounds and objectness scores. This is done by adding two convolutional layer after the shared convolutional layers. The first one converts extracted feature map into a feature vector and the second one generates an objectness score and a regressed bounds for let's assume k region proposals as in figure 4. To generate region proposals, a small convolutional network slides over the shared feature map and each of these sliding windows are mapped to a lower-dimensional feature.

These k region proposals are fed into detection network with the extracted feature map from shared convolutional layers. Then, region of interest pooling extract region proposals from extracted feature map. Predictions are made by using this final feature map. Thus, RPN says the detector network where to look.

Training of Faster R-CNN consists of four steps. First, RPN is trained. Since negative samples dominate positive samples in region proposals, 256 samples are sampled to train RPN by using mini-batches from a single image. Then, in the second step, Fast R-CNN is trained separately by using the region proposals generated by the RPN in step 1. So far, two networks are separate and they do not share convolutional layers. As a step 3, the shared convolutional layers are held fix and RPN is fine-tuned by using detector network. As a final step, the layers of detection network is fine-tuned by holding the shared convolutional layers fixed.




Figure 4: An example of RPN

3.1.2 Mask R-CNN

3.2 One-Stage Detectors

3.2.1 You Only Look Once: Unified, Real-Time Object Detection

You Only Look Once (YOLO) refers to a unified, real-time object detection. YOLO is called unified since it performs simultaneous prediction over all image using its fixed number of bounding boxes in each grid cell. Also, whole detection pipeline is a single network, thus YOLO can be optimized end-to-end directly from given image to its detection performance.

The process of prediction can be examined as follows:

1. The given image is split into $S \times S$ grid.
 - Each grid cell is responsible for prediction of one object. The selection of which cell is responsible is chosen by the center of the object. If an object's center falls into a cell, that cell is responsible to predict that object.
2. Each grid cell has B bounding boxes to make prediction.

- One prediction for each grid cell, independent of number of bounding boxes, causes YOLO to predict limited number of objects when there are multiple object present in a grid cell.
3. Thus, YOLO predicts B bounding boxes and one confidence score for each grid cell. Now, let's get into more detail for these predictions:
- For each of these bounding boxes, five values are predicted: x , y , w , h , and confidence score. The values x and y are corresponding object's center coordinates on given image. The values w and h are corresponding object's width and height. The confidence score refers to detector's confidence whether a object exists in the corresponding bounding box and how accurate the bounding box is.
 - The values x and y are normalized relative to the cell they belong and the values w and h are normalized relative to the given image, so that, all these x , y , w , and h values are between 0 and 1. Also, a class conditional probability for each grid cell is predicted. Thus, each category in possible class set has one class conditional probability in each grid cell.
 - Thus, for $S \times S$ grid and B bounding boxes for each grid cell, predictions are encoded as $S \times S \times (B * 5 + C)$ tensor.

3.2.2 YOLO9000: Better, Faster, Stronger

Before introducing YOLO9000, one has to explain YOLOv2 since YOLOv2 is the improved version of YOLO in both detection performance and speed while keeping real-time speed of detection. YOLO9000 refers to a joint training method on object detection and classification. Using this joint training method, one can detect objects that do not have any labelled detection data. Now, YOLOv2 will be explained and later, YOLO9000 will be explained.

As the name of the article suggested, YOLOv2 will be explained in three different perspective: Better, faster, stronger. While better and faster refer to performance increase of YOLO which is YOLOv2, stronger refers to YOLO9000.

1. Better

- Analyzing the sufferings of YOLO, one can see that YOLO suffers from localization and having a low recall rate. Thus, the aim is to improve recall and localization while maintaining classification accuracy.

(a) Batch Normalization

- Applying batch normalization leads to significant improvement in performance and also it reduces the need for other forms of regularizations. One can remove dropout from the model without overfitting.
- Batch normalization leads to 2% increase in mAP.

(b) High Resolution Classifier

- Training of YOLO is occurred in two phases: first, convolutional layers are trained on ImageNet classification task on image resolution of 224 x 224. Then, the convolutional layers are trained for detection on image resolution of 448 x 448. This causes detector to switch itself for learning detection for new resolution simultaneously.
- In YOLOv2, before detector is switched for learning detection, classifier network is trained for 448 x 448 resolution images for 10 epochs on ImageNet. Thus, a time is provided for network to adjust itself to work better on high resolution.

- High Resolution Classifier leads to 4% increase in mAP.
- (c) Convolutional with Anchor Boxes
 - convolutional with anchor boxes
- (d) Dimension Cluster
 - dimension cluster
- (e) Direct Location Prediction
 - direct location prediction
- (f) Fine-grained Features
 - fine-grained features
- (g) Multi-scale Training
 - The aim of multi-scale training is to make detector more robust to running images of different sizes. Thus, detector chooses a new image dimension size in every 10 epochs during training from a list of sizes {320, 352, ..., 608}.
 - There is a tradeoff between speed and accuracy since detector performs faster when image size is smaller but accuracy is worse when image size is smaller.

2. Faster

The main motivation of YOLO is to design a detector which performs real-time object detection while maintaining the detection accuracy high as possible. To increase detection performance, some features of detector has to be more complex but designing a more complex architecture causes slower detection speed. To create space for more complex features, in YOLOv2, speed of backbone network is increased as following:

- VGG16 requires 30.69 billion floating point operations for a single pass over a single image at 224 x 224 resolution [reference to paper!].
- YOLO framework uses a custom backbone network based on GoogleNet architecture. This network uses 8.52 billion operations, however, it works slightly worse than VGG16.




Figure 5: Network of SSD

- Thus, in YOLOv2, a new classification model is used as backbone network which is called DarkNet19. This network requires 5.58 billion operations.
- Their accuracy on ImageNet as follows:
 - (a) VGG16: 90.0% top-5 accuracy.
 - (b) YOLO: 88% top-5 accuracy.
 - (c) DarkNet19: 91.2% top-5 accuracy.

3. Stronger

Word-tree hierarchy!




Figure 6: Multi-scale feature maps




Figure 7: Default Bounding Boxes

3.2.3 YOLOv3: An Incremental Improvement

3.2.4 SSD: Single Shot MultiBox Detector

Single Shot MultiBox Detector is a single deep neural network which aims to detect objects in real-time. The main improvement in speed comes from eliminating region proposals as Faster R-CNN does. Although eliminating region proposals increases the speed of detection, it reduces the accuracy significantly. Thus, three main improvements are applied to increase accuracy:

1. Multi-Scale Feature Maps for Detection

- As mentioned in section 2.3, handling multi-scale is important feature for object detectors. It can reduce the mAP for detector's performance since objects of different sizes in given image cannot be detected well. To handle multi-scale detection, SSD uses multi-scale feature maps instead of using different sizes of input images. The size of feature maps decreases through the network's architecture. You can see in figure 5 the SSD architecture in which layers are getting smaller from left to right and each of these convolutional layers are connected to the prediction module, so that, multi-scale feature maps can be used in prediction. For larger feature maps, SSD can detect smaller objects and for smaller feature maps, smaller objects can be detected. You can see in figure 6, detecting cat in image is done by 8x8 feature map, although, detecting dog is done by 4x4 feature map.

2. Convolutional Predictors for Detection

- Small-size convolutional filters are used to make prediction for object detection. These convolutional filters are applied on extracted feature maps to compute both localization and class scores. Each filter computes a bounding box and corresponding class scores for each category.

3. Default Bounding Boxes and Aspect Ratios

- SSD divides its feature maps into a grid and each feature map cells are associated with a set of default bounding boxes. These

bounding boxes have fixed position relative to its corresponding grid cell. The aim of using multiple bounding boxes in each grid cell is to detect different-shape objects such as cars and people. You can see in figure 7, there are four different bounding boxes for each grid cells. These bounding boxes are shaped different than the others to increase the chance of detection of different-shaped objects. For example, bounding box 1 in figure 7 can be used to detect cars but bounding box 3 is better option for detecting people.

- For different layers of SSD, default bounding boxes are customized for different resolutions. There are target aspect ratios and corresponding to these aspect ratios, default bounding boxes are calculated.

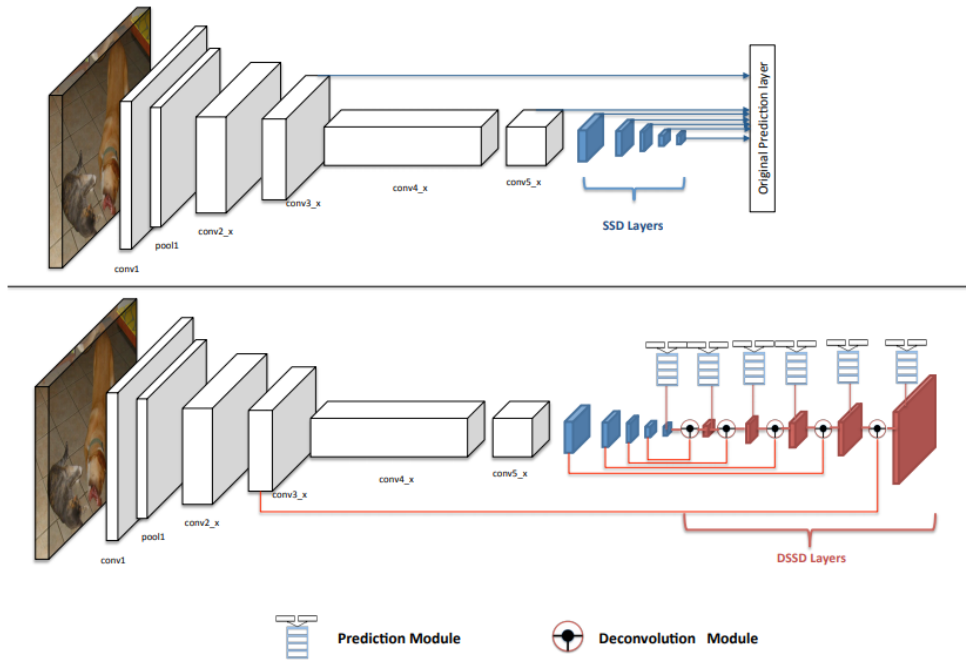


Figure 8: Networks of SSD and DSSD on ResNet

3.2.5 DSSD: Deconvolutional Single Shot Detector

Deconvolutional Single Shot Detector aims to contribute a new approach for object detection. The major changes in DSSD are to change the base network from VGG16 to ResNet101, using prediction modules, and adding deconvolution layers after convolutional layers of SSD. In figure 8, you can see the DSSD layers and the way they are connected with the layer before and corresponding size SSD layer. Also, the change in prediction layer from SSD to DSSD can be seen. So, these changes will be examined as follows:

1. ResNet101

- The aim of changing base network from VGG16 to ResNet101 is to increase accuracy. However, it does not improve accuracy by itself. That's why, the prediction module is used to increase performance.

2. Prediction Module

- Due to the principle of improving sub-network can improve accuracy, original SSD approach for parameters prediction is replaced with a prediction module which consists of a residual block. Thus, using ResNet101 with prediction module performs better than VGG16. The implementation of prediction module can be seen in figure 9.

3. Deconvolution Layers

- The aim of using deconvolutional layers is to include more high-level context in detection, so that, deconvolution module integrates information from both earlier feature maps and earlier deconvolution layers. Instead of using deconvolution module, up-sampling layers would have been used to increase the resolution of feature maps such as in hourglass networks. However, using deconvolution module provides learnable parameters which performs better. You can see the implementation of deconvolution module in figure 10.




Figure 9: Prediction Module




Figure 10: Deconvolution Module

3.2.6 CornerNet: Detecting Objects as Paired Keypoints

CornerNet is a complete new approach to object detection. It detects objects using paired keypoints which means detecting bounding boxes using top-left and bottom-right corners of objects. In addition to this new approach, corner pooling is introduced. Corners are better localized using corner pooling. Using paired keypoints eliminates the need for using anchor boxes.

There are two drawbacks of using anchor boxes:

1. Anchor boxes require a huge set. Since most of the anchor boxes do not overlap with ground-truth bounding boxes, a huge imbalance between positive and negative anchor boxes is created and handling this imbalance slows down training.
2. Anchor boxes introduce many hyperparameters and design choices to make, it may even introduce more if a single network makes separate predictions at multiple resolutions.

Corner pooling is applied by taking the maximum values in two directions (horizontal and vertical) for each channel and adding these two values together.

In CornerNet, a single neural network predicts a heatmap for the top-left corners of all instances of the same object category, a heatmap for all bottom-right corners, and an embedding vector for each detected corner [reference for paper!].




Figure 11: Focal Loss

3.2.7 FPN

3.2.8 Focal Loss and RetinaNet

In general, two-stage detectors have a better accuracy than one-stage detectors. The aim of this article is to find out why this is the case. Two-stage detectors are applied on a sparse set of candidate object locations. However, one-stage detectors are applied over dense sampling of possible object locations. Then, the obtained result is the foreground-background class imbalance in training of dense detectors. The improved solution to this problem is called focal loss (Fig. 11) which is introduced by this article.

Focal loss is the reshaped version of cross-entropy loss. The aim of this change in cross-entropy is to down-weights the loss calculated for well-classified examples. Thus, detector is trained on sparse set of hard examples. Also, class imbalance is handled by focal loss and sampling examples is not required since well-learned examples do not overwhelm loss during training. You can see the change in loss by looking figure 11. When γ equals to zero, focal loss is equivalent to cross-entropy loss. For higher values of γ , the convergence of loss values changes as in the figure 11. The best-performing γ value equals to two according to the article.

RetinaNet is designed to show the effectiveness of Focal Loss. As you can see in figure 12, RetinaNet uses ResNet as its backbone network and

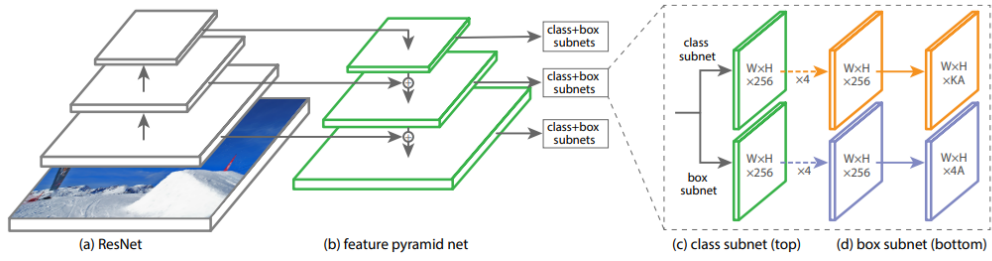


Figure 12: Network of RetinaNet

FPN to obtain a multi-scale convolutional feature pyramid. Then, two sub-networks are used to obtain detection results. One of them is used to classify anchor boxes to obtain classification of objects in anchor boxes and the other one is used to regress bounding boxes from anchors in which the aim is to obtain ground-truth bounding boxes.

3.2.9 EfficientDet: Scalable and Efficient Object Detection

3.2.10 FreeAnchor: Learning to Match Anchors for Visual Object Detection

4 Datasets and Metrics

In this section datasets and performance metrics will be explained.

Datasets are used in training and testing of detectors. Datasets consist of some specific type of objects, in object detection case they are images, and the corresponding ground-truth information about the categories in images that detector performs to detect. So, the following datasets are mainly used in object detection:

- Pascal VOC
- MS COCO
- MOT

Performance metrics are the way to measure the accuracy of detectors. In object detection, improving precision and recall yields to better test performance. The difference lies under the calculation of these metrics for object detection since correctly classifying the object in the image is not enough, detector also correctly localize the object in the image. So, to measure the correctness of each detection, a metric called Intersection over Union is used:

- Intersection over Union (IoU) IoU is the ratio between the intersection of predicted and ground-truth bounding box divided by union of predicted and ground-truth bounding box.
- Correctness of Detection IoU is used as a threshold to identify a detection as positive or negative. The most commonly used IoU threshold is 0.5 which means if IoU value of a detection is bigger than 0.5, that detection is counted as True Positive; otherwise, False Positive.
- The following concepts are used by the metrics:
 - True Positive (TP): A correct detection, $\text{IoU} \geq \text{threshold}$.
 - False Positive (FP): A wrong detection, $\text{IoU} \leq \text{threshold}$.

Stage	Model-Input Resolution	Backbone	mAP	FPS
Two Stage	Faster R-CNN	VGG16	39.3	5
Two Stage	Faster R-CNN + FPN	ResNet101	58.5	6.75
One Stage	YOLOv2	DarkNet19	21.6	40.0
One Stage	YOLOv3-320/416/608	DarkNet53	28.2/31.0/33.0	45.45/34.48/19.60
One Stage	SSD-300/512	VGG16	25.1/28.8	59.0/22.0
One Stage	SSD-513	ResNet101	31.2	8.0
One Stage	DSSD-513	ResNet101	33.2	6.41
One Stage	RetinaNet500/800	ResNet50/101	32.5/37.8	13.88/5.05
One Stage	RetinaNet	ResNet101/X	39.1/40.8	8.19
One Stage	EfficientDet-D3	EfficientNet	44.3	23.81
One Stage	CornerNet(s/m)	Hourglass	40.5/42.1	4.09

Figure 13: Performance table

- False Negative (FN): A ground-truth not detected.
- True Negative (TN): Since there are many possible bounding boxes that do not contain any object, counting corrected misdetection is not necessarily needed.

5 Performance Comparison

Understanding the performance of detectors based on different cases and different features is important knowledge to decide on usage of detectors in real-life problems. Therefore, a table to show the performance of detectors is prepared which can be seen in figure 13.

5.1 Accuracy and Real-Time Applicable

Obtained accuracy and frame rate per second performances of detectors are visible on figure 13. These results are valid on COCO dataset. Since most of the detectors perform well on Pascal VOC dataset and COCO is more hard dataset to obtain better accuracy, only the results on COCO dataset is shown. Also, the aim of this study is to find out which detectors can be applicable in real-time. Therefore, the detectors that are shown in figure 13 are chosen in a specific form in which they are most close to their real-time applicable specifications.

Based on the results in figure 13, YOLOv2, YOLOv3 for 320x 320 image resolution, YOLOv3 for 416x416 image resolution, SSD for both 300x300 and 512x512 image resolution, EfficientDet-D3 can be used in real-time applications since videos require 24 FPS to play without any distortion in human perception. Although, rest of the detectors cannot be applicable in real-time, their accuracy is better than these real-time applicable detectors. Thus, these detectors can be used in different cases than real-time applications, so that, better accuracy in objective can be obtained.

5.2 On COCO

5.3 In Video

5.3.1 MOT17Det

5.4 Quantitative Analysis

5.4.1 Based on Size

5.4.2 Based on Location on Image

5.4.3 Based on Visibility Ratio

5.4.4 Based on Confidence Scores

5.4.5 Based on Combinations

5.5 Qualitative Analysis